

## TML Assignment - SARSA, Expected SARSA and Q Learning

Rehas Mehar Kaur Sachdeva - 201401102

Run on a 4 by 4 grid world with (0,0) as start state, (3,3) as goal state and (1,2) as a bad state. Notation - 0 : Left, 1 : Right, 2 : Up, 3 : Down.

### SARSA

Epsilon = 0.05


Optimal Policy

```
[[ 2.  1.  2.  1.]  
 [ 2.  1.  1.  1.]  
 [ 2.  1.  1.  1.]  
 [ 0.  2.  2.  0.]]
```

Optimal Value functions

```
[[[ 2.68517441e+00  3.67435563e+00  3.08867837e+01  1.28790810e+01]  
 [ 1.41845698e+01  4.13431219e+01 -5.07410469e-02  1.00140419e+01]  
 [-1.82335000e+00 -3.61131971e+01  1.25919132e+01 -1.99000000e+00]  
 [-1.82335000e+00  4.75906891e+01 -8.71507990e+00 -1.95575500e+01]]  
  
[[ -2.26671653e+00 -2.10203830e+00  3.45395743e+01 -2.24905330e+00]  
 [ 1.85542633e+01  5.73219806e+01 -7.00000000e+01  1.64597968e+01]  
 [ 7.73984791e-01  6.77842623e+01 -1.99000000e+00 -1.00000000e+00]  
 [-3.55000000e+01  7.80660069e+01 -1.00000000e+00 -7.00000000e+01]]  
  
[[ 9.59794100e-01 -1.51826748e+00  5.97021704e+01  1.51970297e+01]  
 [ 2.65438447e+01  7.87243700e+01  3.85088000e+01  1.28439728e+01]  
 [-3.81871770e+01  8.87428588e+01  3.52402639e+01 -2.97010000e+00]  
 [-1.00000000e+00  9.61713390e+01  0.00000000e+00  0.00000000e+00]]  
  
[[ 2.89915996e+01  4.64597000e+00 -1.00000000e+00 -1.49500000e+00]  
 [ 4.26441397e+01  4.22169819e+01  8.94773669e+01  2.33431170e+01]  
 [ 6.16804729e+01  6.81159648e+01  9.71796413e+01  7.67198847e+01]  
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00]]]
```

## Reward per episode



Epsilon = 0.2

## Optimal Policy

```
[[ 1.  3.  2.  1.]  
 [ 1.  1.  1.  1.]  
 [ 2.  2.  1.  1.]  
 [ 2.  2.  2.  0.]]
```

## Optimal Value functions


```
[[[ 1.05913734 12.72231838 -0.18762947  2.17831198]  
 [ -2.27628048 -5.17786607 -6.92610812  7.00936106]  
 [ -1.48071775 -36.48520349  4.1860226  -1.99   ]  
 [ -1.495    22.9211307  -1.92626875 -1.99   ]]]
```

```
[[ 5.23429572 30.23934273 24.61642063 10.5466985 ]  
 [ -7.92555031 43.42754466 -42.99891544 -1.83978963]  
 [ -1.        64.88006128 28.10848544 35.20503452]  
 [ 7.8693247  68.75175103 32.73128652 -60.13877812]]
```

```
[[ 12.83641056 14.88686271 50.12586225 13.39634576]
 [ 31.49178651 47.58000115 68.41118569 27.42709646]
 [-35.23093391 88.72288928 77.5492931 49.40030817]
 [ 21.11119901 95.99739027 54.83824483 87.66444648]]
```

```
[[ -2.22438827 -2.62493382 42.50285053 7.93723796]
 [ 36.49980811 32.20981207 73.82058494 -1.99    ]
 [ 41.86620214 67.9556181 96.48201792 59.4535089 ]
 [ 0.         0.         0.         0.        ]]]
```

## Reward per episode



## Effect of Epsilon on Rate of convergence and optimal policy

Takes slightly longer (~40 more iterations) to converge for epsilon=0.2. Epsilon=0.05 seems to converge to better policy, but the final reward they converge to are almost same.

# Expected SARSA

Epsilon = 0.05


## Optimal Policy

```
[[ 1.  1.  2.  1.]  
 [ 1.  1.  1.  1.]  
 [ 1.  1.  3.  1.]  
 [ 2.  2.  2.  0.]]
```

## Optimal Value functions

```
[[[ 4.44956751 25.69825885 -2.13382104  2.35475723]  
 [ -2.49587152 16.51568866 -2.27035961 -2.41292394]  
 [ -2.55439801 -70.         4.19124546 -2.58330734]  
 [ -2.01248152 32.32477917 -2.09420559 -2.70015564]]  
  
[[ 2.06167022 45.97037987  3.39090972  2.80990192]  
 [ -2.01934187 45.02987161 -70.51765665 -1.53820316]  
 [ -1.87870157 26.47260263 -1.01268223 -1.89159493]  
 [ -1.02482657 75.21985193 -1.03139977 -70.04815587]]  
  
[[ 16.77443987 66.1993347  34.8825798  42.15782294]  
 [ -1.02505723 76.79363319 -1.87935171 -2.86802905]  
 [ -70.03562393 -1.0231003 -1.01268508 62.52473399]  
 [ -1.02505723 100.         0.         0.         ]]  
  
[[ 28.8117151  24.26006905 80.65656162 36.91371893]  
 [ 58.57520113 57.13330652 90.25652742 62.23570261]  
 [ 38.15310606 60.95438523 97.26729612 85.62567911]  
 [  0.         0.         0.         0.         ]]]
```

## Reward per episode



Epsilon = 0.2

## Optimal Policy

```
[[ 1.  3.  0.  0.]  
 [ 1.  3.  1.  1.]  
 [ 1.  1.  1.  1.]  
 [ 2.  2.  2.  0.]]
```


## Optimal Value functions

```
[[[ -1.34080785  6.7896073 -3.26989679 -1.40539162]  
 [ -3.28717199 -4.39746642 -4.51817538 -1.06818724]  
 [ -1.02475   -70.0495   -1.88716135 -1.22364968]  
 [ -1.95020693 -3.23995545 -2.49041859 -3.30300988]]  
  
[[ -1.75848928  21.10152633 -4.56811208  7.09622304]  
 [ -1.61216682 -36.17222375 -60.52194081  5.17984448]  
 [ -2.78937302  23.34716558 -1.10145025 -4.21576581]  
 [ -1.58803254  48.92438709 -1.12871114 -70.32503929]]
```

```
[[ 13.87728692  39.92695972  34.75632287   9.95958937]
 [ -5.9085166   63.4464957  -4.99500515  11.28772105]
 [ -47.96168312  80.25907572  41.02427488  -1.41434934]
 [ 14.44276596 100.      -1.0495    0.    ]]
```

```
[[ 14.50711651  15.71168546  61.24020121  27.95117894]
 [ 29.96042575  41.04120089  84.12986366  29.57971726]
 [ 50.28974817  66.34947216  96.9731518   63.52371588]
 [ 0.         0.         0.         0.    ]]
```

## Reward per episode



## Effect of Epsilon on Rate of convergence and optimal policy

Takes about 35 more iterations for epsilon=0.05 to converge but it converges to a significantly better policy than for epsilon=0.2, with higher reward at convergence.

# Q Learning

Epsilon = 0.05


Optimal Policy

```
[[ 1.  1.  3.  0.]  
 [ 2.  1.  1.  1.]  
 [ 2.  2.  2.  1.]  
 [ 1.  2.  2.  0.]]
```

Optimal Value functions

```
[[[ 3.79105358 29.68351057 6.92033283 0.76459648]  
  [-2.06490849 38.11166321 -1.99      12.59404944]  
  [-1.82335   -70.      -1.66     -0.51597695]  
  [-1.33      -1.740025  -1.495    -1.76589147]]  
  
[[ -1.33383856  3.70834858 49.75447866 13.97230549]  
 [ 15.06349965 68.40910654 -17.33234627 16.58639499]  
 [-1.39041117 78.94074807 -1.      -1.      ]  
 [-1.495     -1.      -1.      -70.     ]]  
  
[[ -1.79068     1.53607982 53.51924235 -2.01153745]  
 [ 22.90671556 14.35489668 85.31616926 6.46395398]  
 [-30.49165888 50.70966156 94.88765344 40.57804717]  
 [-1.      99.4724454  0.      90.46687647]]  
  
[[ -1.99      -1.82335   -1.95535   -2.06879163]  
 [-1.6175125  -1.82335   70.17828518 -1.8217   ]  
 [-1.495     -1.495     94.33850923 -1.99     ]  
 [ 0.      0.      0.      0.      ]]]
```

## Reward per episode



Epsilon = 0.2

## Optimal Policy

```
[[ 1.  1.  3.  1.]  
 [ 2.  1.  1.  1.]  
 [ 2.  1.  2.  1.]  
 [ 2.  2.  2.  0.]]
```

## Optimal Value functions


```
[[[ 5.61071599e+00  3.01930433e+01  7.87716373e+00  6.38436679e+00]  
 [ -2.22338008e+00  3.34282753e+01 -2.23502500e+00  7.63119049e+00]  
 [ -2.53176071e+00 -3.55000000e+01 -7.00000000e+01  5.20698362e+00]  
 [ -1.49500000e+00 -1.00000000e+00 -1.00000000e+00 -2.16177138e+00]]  
  
[[ 1.74484924e+01  3.28710250e+01  4.68344693e+01  2.49311646e+01]  
 [ 9.42524509e+00  6.49351582e+01 -2.09308815e+01  2.97049233e+01]  
 [ -1.55711675e+00  6.10135830e+01  2.07562674e+01  2.33080930e+01]  
 [ -1.49500000e+00  4.73242402e+01 -1.00000000e+00 -4.53717261e+01]]
```



```
[[ 1.21292460e+01  3.59715993e+00  6.85472037e+01  4.64727416e-02]
 [ 4.64600849e+01  8.51033801e+01  5.60727565e+01  4.42514207e+01]
 [-4.25246829e+01  4.84248517e+01  8.36574880e+01  2.46425181e+01]
 [-1.00000000e+00  9.90485066e+01  7.22401453e+01  6.32551431e+01]]
```

```
[[ 8.52266185e+00  1.42327500e+01  8.67142599e+01  9.07908541e+00]
 [ 5.23213686e+01  7.93040687e+01  9.55689187e+01  5.69686765e+01]
 [ 6.61470410e+01  7.60122809e+01  9.91638477e+01  9.30536751e+01]
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00]]]
```

## Reward per episode



## Effect of Epsilon on Rate of convergence and optimal policy

Both have very similar rate of convergence and very similar optimal policy.